


# COMBINING MODELS



# Material

- Mehta et al notebooks:
  - 08 Bagging
  - 09 Random Forests
  - 10 XGBoost
- Our XGBoost notebooks



# Combining models

- “wisdom of the crowds”

(if no correlated deficiencies)

One of the most powerful and widely-applied ideas in modern machine learning is the use of ensemble methods that combine predictions from multiple, often weak, statistical models to improve predictive performance (Diet-

- Many weak models need less assumptions than a single complicated model
- Today: widely applied techniques


# Boosting Algorithms: AdaBoost, Gradient Boosting and XGBoost



Rohith Gandhi [Follow](#)

May 6, 2018 · 5 min read

<https://hackernoon.com/boosting-algorithms-adaboost-gradient-boosting-and-xgboost-f74991cad38c>




Here, let A, B and C be different classifiers. Their area A represents where the classifier A misclassifies (goes wrong) and area B represents where the classifier B misclassifies and area C represents where the classifier C misclassifies. Since there is no correlation between the errors of each classifier, combining them and using a technique of democratic voting to classify each object, this family of classifiers will never go wrong.

# Correlated vs uncorrelated models

- Important that combined models are uncorrelated (See the review)
- **Correlated** models → not so better variance, possibly higher bias
- M **uncorrelated** models → much reduced variance ( $\sim 1/M$ ), same bias of single model


Is there a combination of linear classifiers that predicts the color of these data?



# 1) bagging








Notebook 8: Bagging a simple binary classifier

- subsample data randomly, many times: **bootstrapping**



- apply a **simple** model to each subset
- Combine outputs of models with **majority rule**
- <https://datasciencelab.wordpress.com/2014/01/10/machine-learning-classics-the-perceptron/>








## 2) decision trees












# 2) decision trees

Changing order



## 2) decision trees



# 3) boosting (AdaBoost)

- Form aggregate classifier iteratively
- works on improving the areas where the base learner fails
- <https://web.stanford.edu/~hastie/Papers/AdditiveLogisticRegression/alr.pdf>

---

## Discrete AdaBoost [Freund and Schapire (1996b)]

1. Start with weights  $w_i = 1/N, i = 1, \dots, N$ .
  2. Repeat for  $m = 1, 2, \dots, M$ :
    - (a) Fit the classifier  $f_m(x) \in \{-1, 1\}$  using weights  $w_i$  on the training data.
    - (b) Compute  $\text{err}_m = E_w[1_{(y \neq f_m(x))}]$ ,  $c_m = \log((1 - \text{err}_m)/\text{err}_m)$ .
    - (c) Set  $w_i \leftarrow w_i \exp[c_m 1_{(y_i \neq f_m(x_i))}]$ ,  $i = 1, 2, \dots, N$ , and renormalize so that  $\sum_i w_i = 1$ .
  3. Output the classifier  $\text{sign}[\sum_{m=1}^M c_m f_m(x)]$ .
-

# 3) boosting (AdaBoost)

---

## Real AdaBoost

1. Start with weights  $w_i = 1/N$ ,  $i = 1, 2, \dots, N$ .
  2. Repeat for  $m = 1, 2, \dots, M$ :
    - (a) Fit the classifier to obtain a class probability estimate  $p_m(x) = \hat{P}_w(y = 1|x) \in [0, 1]$ , using weights  $w_i$  on the training data.
    - (b) Set  $f_m(x) \leftarrow \frac{1}{2} \log p_m(x)/(1 - p_m(x)) \in \mathbb{R}$ .
    - (c) Set  $w_i \leftarrow w_i \exp[-y_i f_m(x_i)]$ ,  $i = 1, 2, \dots, N$ , and renormalize so that  $\sum_i w_i = 1$ .
  3. Output the classifier  $\text{sign}[\sum_{m=1}^M f_m(x)]$ .
- 

ALGORITHM 2. *The Real AdaBoost algorithm uses class probability estimates  $p_m(x)$  to construct real-valued contributions  $f_m(x)$ .*

\* Drawback: easily defeated by noisy data,  
the algorithm tries to fit every point perfectly, hence outliers are a problem